

Raft3D - 3D Prints Management using Raft Consensus Algorithm

The Raft3D project is a distributed system designed to manage 3D printing resources (printers, filaments, and print jobs) using a simplified mock Raft consensus algorithm. Implemented as a Dockerized cluster with three nodes (raft3d-node1, raft3d-node2, raft3d-node3), it demonstrates fault tolerance, leader election, RESTful API functionality, and monitoring through Prometheus metrics. Below is a comprehensive description of what was implemented and developed, focusing on the system's features, components, and capabilities.

Project Overview

Purpose: Raft3D provides a fault-tolerant system for managing 3D printing resources in a distributed environment. It uses a mock Raft algorithm (MockRaft) to ensure a single leader coordinates operations, exposes APIs for resource management, and provides metrics for monitoring system health.

Key Objectives:

1. Deploy a three-node cluster using Docker Compose.
2. Implement leader election to maintain a single leader at a time.
3. Develop RESTful APIs for managing printers, filaments, and print jobs.
4. Enable fault tolerance by electing a new leader when a node fails.
5. Expose Prometheus metrics to monitor leader status, requests, and snapshots.
6. Support periodic state snapshots for persistence.

Core Components:

- **Dockerized Cluster:** Three nodes running as Docker containers with shared storage.
- **Mock Raft Algorithm:** Simplified Raft implementation for leader election using a shared file.

- **FastAPI Application:** RESTful APIs for managing printing resources.
- **Prometheus Metrics:** Metrics for leader status, HTTP requests, and snapshots.
- **Snapshot Manager:** Periodic state snapshots stored in a shared directory.

Implemented Features and Components

1. Dockerized Cluster

- **Docker Compose Configuration:**
 - Defined three services (raft3d-node1, raft3d-node2, raft3d-node3) in docker-compose.yml.
 - Configured environment variables: NODE_ID (node identifier), RAFT_PORT (Raft communication), HTTP_PORT (API endpoint), RAFT_DIR (data directory), and CLUSTER (peer list).
 - Mapped external ports 8080-8082 to internal HTTP port 8080 and 9090-9092 to internal Raft port 9090.
 - Used a single shared volume (raft-data) mapped to /raft/data for all nodes to store leader.txt and snapshots.
- **Deployment:**
 - Built and started the cluster with `sudo docker-compose up --build -d`.
 - Verified container status with `sudo docker ps`, confirming container names (e.g., raft3d_raft3d-node1_1) and port mappings.
- **Capabilities:**
 - Runs three independent nodes communicating over a bridge network (raft3d-net).
 - Ensures shared storage for consistent leader coordination and state persistence.

2. Mock Raft Implementation

- **MockRaft Class** (app/raft/mock_raft.py):
 - Implemented a simplified Raft algorithm to manage leader election and state application.
 - Used a shared leader.txt file in /raft/data to track the current leader's ID.
 - Provided methods:
 - `start`: Initiates a background thread for leader election.
 - `stop`: Stops the node and resigns leadership if applicable.

- `become_leader`: Writes the node's ID to `leader.txt` after verifying no other leader exists.
- `resign_leader`: Removes `leader.txt` if the node is the leader and updates metrics.
- `run`: Periodically checks for a leader, removes stale `leader.txt` (>10 seconds old), and attempts election with 50% probability.
- `apply`: Applies log entries to a finite state machine (FSM) for state updates.
- Integrated with `app/raft/store.py` (`RaftNode`) to manage state transitions.
- **Leader Election:**
 - Nodes check `leader.txt` to detect the current leader.
 - If no leader exists or `leader.txt` is stale, nodes attempt election by writing their ID to `leader.txt`.
 - Elections occur with a 50% probability to balance frequency and stability.
- **Fault Tolerance:**
 - Detects stale `leader.txt` files (older than 10 seconds) to recover from leader crashes.
 - Ensures only one leader exists at a time through shared file coordination.
- **Capabilities:**
 - Maintains a single leader to coordinate operations.
 - Supports recovery from node failures by electing a new leader.
 - Applies log entries to ensure consistent state updates.

3. RESTful APIs

- **FastAPI Application** (`app/main.py`, `app/api/handlers.py`):
 - Developed a FastAPI server for each node to handle API requests.
 - Implemented endpoints under `/api/v1`:
 - **Printers:**
 - `POST /printers`: Create a printer (e.g., `{"id": "p1", "company": "Creality", "model": "Ender 3"}`).
 - `GET /printers`: Retrieve all printers.

- **Filaments:**
 - POST /filaments: Create a filament (e.g., {"id": "f1", "type": "PLA", "color": "Blue", "total_weight_in_grams": 1000}).
 - GET /filaments: Retrieve all filaments.
- **Print Jobs:**
 - POST /print_jobs: Create a print job (e.g., {"id": "j1", "printer_id": "p1", "filament_id": "f1", "filepath": "prints/sword/hilt.gcode"}).
 - POST /print_jobs/{id}/status: Update print job status (e.g., Queued, Running, Done).
- Integrated with RaftNode to apply API requests to the Raft log for consistency.
- **Capabilities:**
 - Enables creation, retrieval, and updating of 3D printing resources.
 - Supports state updates (e.g., filament weight reduction after print jobs).
 - Operates on any node, with the leader handling state changes.

4. Prometheus Metrics

- **Metrics Setup** (app/monitoring/metrics.py):
 - Defined Prometheus metrics:
 - raft3d_is_leader: Gauge indicating leader status (1 for leader, 0 for follower) per node.
 - raft3d_requests_total: Counter for HTTP requests, labeled by endpoint and method.
 - raft3d_snapshots_total: Counter for snapshots taken.
 - Python process metrics (e.g., process_cpu_seconds_total, process_virtual_memory_bytes).
 - Exposed metrics via /metrics endpoint in FastAPI.
- **Capabilities:**
 - Monitors leader status across nodes.
 - Tracks HTTP request volume and snapshot activity.
 - Provides system-level insights (CPU, memory, file descriptors) for performance monitoring.

5. Fault Tolerance

- **Implementation:**

- Designed the system to handle node failures by electing a new leader.
- Used stale leader.txt detection to identify crashed leaders and trigger elections.
- Provided commands to test fault tolerance:
 - Stop the leader node (e.g., `sudo docker stop raft3d_raft3d-node1_1`).
 - Verify new leader election via logs and metrics.
 - Perform API operations on a follower node to confirm continued functionality.
 - Restart the stopped node to rejoin the cluster.

- **Capabilities:**

- Ensures system availability by electing a new leader after a failure.
- Maintains API functionality through the new leader or followers.
- Supports node recovery by allowing stopped nodes to rejoin.

6. 6. Snapshots

- **Snapshot Manager:**

- Implemented a SnapshotManager to periodically save state snapshots in `/raft/data`.
- Integrated with `raft3d_snapshots_total` metric to track snapshot counts.
- Provided command to verify snapshots: `bash`
`sudo docker exec raft3d_raft3d-node2_1 ls /raft/data`

- **Capabilities:**

- Persists system state for recovery or inspection.
- Stores snapshots in the shared raft-data volume for consistency.

The Raft3D project successfully implements a distributed 3D printing management system with a mock Raft algorithm, Dockerized deployment, RESTful APIs, Prometheus metrics, fault tolerance, and snapshot persistence. The system supports leader election, resource management, and monitoring, with a complete set of commands to demonstrate its functionality. It serves as a practical example of distributed systems concepts, with robust features for operation, recovery, and debugging, ready to showcase fault-tolerant behavior and API-driven resource management.

Demonstration Commands

The following commands were developed to showcase the system's functionality:

1. Cluster Setup:

```
sudo docker-compose up --build -d

sudo docker ps

sudo docker-compose logs | grep "elected leader"
```

2. Metrics Verification:

```
curl http://localhost:8080/metrics | grep
raft3d_is_leader

curl http://localhost:8081/metrics | grep
raft3d_is_leader

curl http://localhost:8082/metrics | grep
raft3d_is_leader
```

3. API Interactions:

```
curl -X POST http://localhost:8080/api/v1/printers -H
"Content-Type: application/json" -d '{"id": "p1",
"company": "Creality", "model": "Ender 3"}
```

```
curl -X POST http://localhost:8080/api/v1/filaments -H
  "Content-Type: application/json" -d '{"id": "f1",
  "type": "PLA", "color": "Blue",
  "total_weight_in_grams": 1000,
  "remaining_weight_in_grams": 1000}'

curl -X POST http://localhost:8080/api/v1/print\_jobs -H
  "Content-Type: application/json" -d '{"id": "j1",
  "printer_id": "p1", "filament_id": "f1",
  "filepath": "prints/sword/hilt.gcode",
  "print_weight_in_grams": 100, "status": "Queued"}'

curl -X POST
http://localhost:8080/api/v1/print\_jobs/j1/status?status=Running

curl -X POST
http://localhost:8080/api/v1/print\_jobs/j1/status?status=Done

curl http://localhost:8080/api/v1/filaments
```

4. Fault Tolerance:

Check Current Leader:

```
curl http://localhost:8080/metrics | grep
raft3d_is_leader

curl http://localhost:8081/metrics | grep
raft3d_is_leader

curl http://localhost:8082/metrics | grep
raft3d_is_leader
```

Stop Leader (assume node1 is leader, check node1 docker container id):

```
sudo docker stop raft3d_raft3d-node1_1
```

Verify Leader Election:

```
curl http://localhost:8081/metrics | grep  
raft3d_is_leader
```

```
curl http://localhost:8082/metrics | grep  
raft3d_is_leader
```

5. Snapshots:

```
sudo docker exec raft3d_raft3d-node2_1 ls /raft/data
```

Capabilities Summary

The Raft3D system provides:

- **Distributed Cluster:** Three-node setup with shared storage for consistent state management.
- **Leader Election:** Mock Raft ensures a single leader using leader.txt coordination.
- **Fault Tolerance:** Recovers from node failures by electing a new leader and continuing operations.
- **RESTful APIs:** Manages 3D printing resources with create, retrieve, and update operations.
- **Monitoring:** Exposes Prometheus metrics for leader status, requests, snapshots, and system health.
- **Persistence:** Saves state snapshots for recovery and inspection.